

LoanTap Case Study

Gautam Naik (gautamnaik1994@gmail.com)

Github link (<https://github.com/gautamnaik1994/LoanTap-ML-CaseStudy>)

About LoanTap

LoanTap is an online platform committed to delivering customized loan products to millennials. They innovate in an otherwise dull loan segment, to deliver instant, flexible loans on consumer friendly terms to salaried professionals and businessmen.

The data science team at LoanTap is building an underwriting layer to determine the creditworthiness of MSMEs as well as individuals. This case study will focus on the underwriting process behind Personal Loan

Business Problem

Loantap aims to develop a machine-learning model to assess whether an individual should qualify for a loan. As data scientists, our task is to analyze a person's attributes and decide whether they should receive a credit line by creating a predictive model. Additionally, we need to provide recommendations and actionable insights.

Metric

- 1. ROC AUC
- 2. Precision
- 3. Recall
- 4. F1 Score

Data Features

Feature	Description
loan_amnt	The listed amount of the loan applied for by the borrower. If at some point in time, the credit department reduces the loan amount, then it will be reflected in this value
term	The number of payments on the loan. Values are in months and can be either 36 or 60
int_rate	Interest Rate on the loan
installment	The monthly payment owed by the borrower if the loan originates
grade	LoanTap assigned loan grade
sub_grade	LoanTap assigned loan subgrade
emp_title	The job title supplied by the Borrower when applying for the loan
emp_length	Employment length in years. Possible values are between 0 and 10 where 0 means less than one year and 10 means ten or more years
home_ownership	The home ownership status provided by the borrower during registration or obtained from the credit report
annual_inc	The self-reported annual income provided by the borrower during registration
verification_status	Indicates if income was verified by LoanTap, not verified, or if the income source was verified
issue_d	The month which the loan was funded
loan_status	Current status of the loan - Target Variable
purpose	A category provided by the borrower for the loan request
title	The loan title provided by the borrower
dti	A ratio calculated using the borrower's total monthly debt payments on the total debt obligations, excluding mortgage and the requested LoanTap loan, divided by the borrower's self-reported monthly income
earliest_cr_line	The month the borrower's earliest reported credit line was opened
open_acc	The number of open credit lines in the borrower's credit file
pub_rec	Number of derogatory public records
revol_bal	Total credit revolving balance
revol_util	Revolving line utilization rate, or the amount of credit the borrower is using relative to all available revolving credit
total_acc	The total number of credit lines currently in the borrower's credit file
initial_list_status	The initial listing status of the loan, Possible values are – W, F
application_type	Indicates whether the loan is an individual application or a joint application with two co-borrowers
mort_acc	Number of mortgage accounts
pub_rec_bankruptcies	Number of public record bankruptcies
Address	Address of the individual

```
In [ ]: from pyspark.sql import SparkSession
from pyspark.sql.functions import split,count,lower, regexp_extract
import pyspark.sql.functions as sf
import matplotlib.pyplot as plt
from pyspark.sql.types import StringType, ArrayType, StructField, StructType, IntegerType, FloatType, DoubleType
from pyspark.ml.feature import StringIndexer, VectorAssembler
# , MinMaxScaler, StandardScaler
import pandas as pd
import seaborn as sns
sns.set_style("whitegrid")
pd.set_option('display.max_columns', None)
from scipy.stats import chi2_contingency
import numpy as np
from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = "all"

from sklearnx import patch_sklearn
patch_sklearn()
from sklearn.preprocessing import StandardScaler, MinMaxScaler, OneHotEncoder, OrdinalEncoder, TargetEncoder
from sklearn.model_selection import train_test_split
from statsmodels.stats.outliers_influence import variance_inflation_factor
# from matplotlib_inline.backend_inline import set_matplotlib_formats
# set_matplotlib_formats('svg')
```

Intel(R) Extension for Scikit-learn* enabled (<https://github.com/intel/scikit-learn-intelx>)

```
In [ ]: # Create a SparkSession
spark = SparkSession.builder \
    .appName("LoanTap") \
    .config("spark.sql.debug.maxToStringFields", 1000) \
    .config("spark.sql.execution.arrow.pyspark.enabled", "true") \
    .config("spark.sql.shuffle.partitions", 1) \
    .config("spark.network.timeout", "120s") \
    .config("spark.executor.heartbeatInterval", "10s") \
    .getOrCreate()
```

Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
24/07/31 18:28:25 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

```
In [ ]: df = spark.read \
    .option("header", "true") \
    .option("inferSchema", "true") \
    .option("multiline", "true") \
    .option("escape", "\\") \
    .csv("./logistic_regression.csv")
df.cache();
```

```
In [ ]: invalid_address_count = 0
def split_address(address):

    address = address.split("\n")[1]
    city,state,zipcode = "Unknown", "Unknown", "Unknown"
    try:
        if(", " in address):
            # 1726 Cooper Passage Suite 129\nNorth Deniseberg, DE 30723
            # state = DE
            # city = North Deniseberg
            # zip = 30723
            city = address.split(",")[0]
            state = address.split(",")[1].split(" ")[1]
            zipcode = address.split(",")[1].split(" ")[2]
        else:
            # USCGC Tran\nFP0 AP 22690
            city, state, zipcode = address.split(" ")
    except:
        global invalid_address_count
        invalid_address_count += 1
    return (city, state, zipcode)

# split_address_udf = sf.udf(split_address, ArrayType(StringType()))

split_address_udf = sf.udf(split_address, StructType([
    StructField("city", StringType(), True),
    StructField("state", StringType(), True),
    StructField("zipcode", StringType(), True)
]))
```

```
In [ ]: to_drop=[]
```

```
In [ ]: df.printSchema()
```

```
root
 |-- loan_amnt: double (nullable = true)
 |-- term: string (nullable = true)
 |-- int_rate: double (nullable = true)
 |-- installment: double (nullable = true)
 |-- grade: string (nullable = true)
 |-- sub_grade: string (nullable = true)
 |-- emp_title: string (nullable = true)
 |-- emp_length: string (nullable = true)
 |-- home_ownership: string (nullable = true)
 |-- annual_inc: double (nullable = true)
 |-- verification_status: string (nullable = true)
 |-- issue_d: string (nullable = true)
 |-- loan_status: string (nullable = true)
 |-- purpose: string (nullable = true)
 |-- title: string (nullable = true)
 |-- dti: double (nullable = true)
 |-- earliest_cr_line: string (nullable = true)
 |-- open_acc: double (nullable = true)
 |-- pub_rec: double (nullable = true)
 |-- revol_bal: double (nullable = true)
 |-- revol_util: double (nullable = true)
 |-- total_acc: double (nullable = true)
 |-- initial_list_status: string (nullable = true)
 |-- application_type: string (nullable = true)
 |-- mort_acc: double (nullable = true)
 |-- pub_rec_bankruptcies: double (nullable = true)
 |-- address: string (nullable = true)
```

```
In [ ]: df.limit(5).toPandas()
```

	loan_amnt	term	int_rate	installment	grade	sub_grade	emp_title	emp_length	home_ownership	annual_inc	verification_status	issue_d	loan_status	purpose	title	dti	earliest_cr
0	10000.0	36 months	11.44	329.48	B	B4	Marketing	10+ years	RENT	117000.0	Not Verified	Jan-2015	Fully Paid	vacation	Vacation	26.24	Jun-
1	8000.0	36 months	11.99	265.68	B	B5	Credit analyst	4 years	MORTGAGE	65000.0	Not Verified	Jan-2015	Fully Paid	debt_consolidation	Debt consolidation	22.05	Jul-
2	15600.0	36 months	10.49	506.97	B	B3	Statistician	< 1 year	RENT	43057.0	Source Verified	Jan-2015	Fully Paid	credit_card	Credit card refinancing	12.79	Aug-
3	7200.0	36 months	6.49	220.65	A	A2	Client Advocate	6 years	RENT	54000.0	Not Verified	Nov-2014	Fully Paid	credit_card	Credit card refinancing	2.60	Sep-
4	24375.0	60 months	17.27	609.33	C	C5	Destiny Management Inc.	9 years	MORTGAGE	55000.0	Verified	Apr-2013	Charged Off	credit_card	Credit Card Refinance	33.95	Mar-

Data Cleaning

Null values

```
In [ ]: for col in df.columns:
        print(col, " : ", df.filter(df[col].isNull()).count())
```

```
loan_amnt : 0
term : 0
int_rate : 0
installment : 0
grade : 0
sub_grade : 0
emp_title : 22927
emp_length : 18301
home_ownership : 0
annual_inc : 0
verification_status : 0
issue_d : 0
loan_status : 0
purpose : 0
title : 1755
dti : 0
earliest_cr_line : 0
open_acc : 0
pub_rec : 0
revol_bal : 0
revol_util : 276
total_acc : 0
initial_list_status : 0
application_type : 0
mort_acc : 37795
pub_rec_bankruptcies : 535
address : 0
```

```
In [ ]: df.select("revol_util").summary().show();
```

summary	revol_util
count	395754
mean	53.79174863677853
stddev	24.452193062711693
min	0.0
25%	35.8
50%	54.8
75%	72.9
max	892.3

In []: df.select("mort_acc").summary().show();

summary	mort_acc
count	358235
mean	1.8139908160844138
stddev	2.1479304671233352
min	0.0
25%	0.0
50%	1.0
75%	3.0
max	34.0

In []: df.groupby('home_ownership').count().sort("count", ascending=False).show()

home_ownership	count
MORTGAGE	198348
RENT	159790
OWN	37746
OTHER	112
NONE	31
ANY	3

Observations

- Majority of the user have mortgage type home ownership

In []: df = df.withColumn(
 "mort_acc",
 sf.when(
 (sf.col("home_ownership") == "MORTGAGE") & (sf.col("mort_acc").isNull()),
 1
).otherwise(sf.col("mort_acc"))
)
df=df.fillna(0, subset=["mort_acc", "pub_rec_bankruptcies", "revol_util"])
df = df.drop("title")

df=df.withColumn("split_address", split_address_udf("address"))

In []: df.select("split_address").show(20, False)

[Stage 99:> (0 + 1) / 1]

split_address
{Mendozaberg, OK, 22690}
{Loganmouth, SD, 05113}
{New Sabrina, WV, 05113}
{Delacruzside, MA, 00813}
{Greggshire, VA, 11650}
{North Deniseberg, DE, 30723}
{East Stephanie, TX, 22690}
{FPO, AE, 30723}
{FPO, AP, 22690}
{Mauricestad, VA, 00813}
{Bartlettfort, NM, 00813}
{South Matthew, MS, 00813}
{West Beckyfort, MS, 70466}
{Shellychester, OR, 29597}
{Lake Andrew, NH, 29597}
{Stevenfort, HI, 30723}
{West Aprilborough, PA, 00813}
{Cummingsshire, NH, 30723}
{Port Kirstenborough, CO, 70466}
{DPO, AE, 05113}

only showing top 20 rows

In []: df = df.withColumn("city", sf.col("split_address.city")) \
 .withColumn("state", sf.col("split_address.state")) \
 .withColumn("zipcode", sf.col("split_address.zipcode")) \
 .drop("split_address")
df=df.drop("address")

In []: df.groupBy('loan_status').count().show()

loan_status	count
Fully Paid	318357
Charged Off	77673

In []: df.groupBy("term").count().show()

term	count
36 months	302005
60 months	94025

In []: split_col = split(df['term'], " ", -1)
df= df.withColumn('term', split_col.getItem(1))

In []: df.groupBy('emp_title').count().sort('count', ascending=False).show()

emp_title	count
NULL	22927
Teacher	4389
Manager	4250
Registered Nurse	1856
RN	1846
Supervisor	1830
Sales	1638
Project Manager	1505
Owner	1410
Driver	1339
Office Manager	1218
manager	1145
Director	1089
General Manager	1074
Engineer	995
teacher	962
driver	882
Vice President	857
Operations Manager	763
Administrative As...	756

only showing top 20 rows

```
In [ ]: df = df.withColumn("emp_title", lower(df["emp_title"]))
# replace Null values with 'unknown'
df = df.fillna('unknown', subset=['emp_title'])
```

```
In [ ]: df.groupBy('emp_title').count().sort('count', ascending=False).show()
```

[Stage 109:> (0 + 1) / 1]

emp_title	count
unknown	22927
manager	5637
teacher	5430
registered nurse	2627
supervisor	2591
sales	2382
driver	2306
owner	2201
rn	2074
project manager	1776
office manager	1638
general manager	1461
truck driver	1288
director	1192
engineer	1188
police officer	1041
vice president	962
operations manager	961
sales manager	961
store manager	941

only showing top 20 rows

```
In [ ]: df.groupBy('emp_length').count().sort('count', ascending=False).show()
```

emp_length	count
10+ years	126041
2 years	35827
< 1 year	31725
3 years	31665
5 years	26495
1 year	25882
4 years	23952
6 years	20841
7 years	20819
8 years	19168
NULL	18301
9 years	15314

```
In [ ]: df = df.withColumn('emp_length', regexp_extract(df['emp_length'], r'(\d+)', 1))

# df.groupBy('grade').avg('emp_length').show()
avg_emp_length=df.groupBy("grade").agg(sf.avg("emp_length").cast("int").alias("avg_emp_length"))

df=df.join(avg_emp_length, "grade", )

df = df.withColumn(
    "emp_length",
    sf.when(
        sf.col("emp_length").isNull(),
        sf.col("avg_emp_length")
    ).otherwise(sf.col("emp_length"))
)

df = df.withColumn("emp_length", df["emp_length"].cast(IntegerType()))
df=df.drop("avg_emp_length")
```

```
In [ ]: df.groupBy("purpose").count().sort('count', ascending=False).show()
```

purpose	count
debt_consolidation	234507
credit_card	83019
home_improvement	24030
other	21185
major_purchase	8790
small_business	5701
car	4697
medical	4196
moving	2854
vacation	2452
house	2201
wedding	1812
renewable_energy	329
educational	257

```
In [ ]: df.groupBy("grade").count().sort('count', ascending=False).show()
```


grade	count
B	116018
C	105987
A	64187
D	63524
E	31488
F	11772
G	3054

```
In [ ]: df.groupby('sub_grade').count().sort('count', ascending=False).show();
```

sub_grade	count
B3	26655
B4	25601
C1	23662
C2	22580
B2	22495
B5	22085
C3	21221
C4	20280
B1	19182
A5	18526
C5	18244
D1	15993
A4	15789
D2	13951
D3	12223
D4	11657
A3	10576
A1	9729
D5	9700
A2	9567

only showing top 20 rows

```
In [ ]: # convert issue_d having format of Jan-2015 to date
df = df.withColumn("issue_d", sf.to_date(df["issue_d"], "MMM-yyyy"))
df = df.withColumn("earliest_cr_line", sf.to_date(df["issue_d"], "MMM-yyyy"))
# extract year from issue_d
df = df.withColumn("issue_year", sf.year(df["issue_d"]))
df = df.withColumn("earliest_cr_line_year", sf.year(df["earliest_cr_line"]))
# extract month in form of integer from issue_d
df = df.withColumn("issue_month", sf.month(df["issue_d"]))
df = df.withColumn("earliest_cr_line_month", sf.month(df["earliest_cr_line"]))
```

```
In [ ]: # # convert multiple colums to int
# df = df.withColumn("earliest_cr_line_year", df["issue_month"].cast(IntegerType()))
# df = df.withColumn("issue_year", df["issue_year"].cast(IntegerType()))
```

```
In [ ]: to_drop.extend(['issue_d', 'earliest_cr_line'])
```

```
In [ ]: # drop duplicates
df = df.dropDuplicates()
```

```
In [ ]: float_cols = [col for col in df.columns if isinstance(df.schema[col].dataType, (DoubleType, FloatType))]
float_cols
```

```
Out[ ]: ['loan_amnt',
'int_rate',
'installment',
'annual_inc',
'dti',
'open_acc',
'pub_rec',
'revol_bal',
'revol_util',
'total_acc',
'mort_acc',
'pub_rec_bankruptcies']
```

```
In [ ]: # convert to int
df=df.withColumn("annual_inc", df["annual_inc"].cast(IntegerType()))
df=df.withColumn("open_acc", df["open_acc"].cast(IntegerType()))
df=df.withColumn("pub_rec", df["pub_rec"].cast(IntegerType()))
df=df.withColumn("total_acc", df["total_acc"].cast(IntegerType()))
df=df.withColumn("mort_acc", df["mort_acc"].cast(IntegerType()))
df=df.withColumn("pub_rec", df["pub_rec"].cast(IntegerType()))
df=df.withColumn("pub_rec_bankruptcies", df["pub_rec_bankruptcies"].cast(IntegerType()))
df=df.withColumn("revol_bal", df["revol_bal"].cast(IntegerType()))
```

```
In [ ]: for col in df.columns:
print(col," : ", df.filter(df[col].isNull()).count())
```

grade : 0
loan_amnt : 0
term : 0
int_rate : 0
installment : 0
sub_grade : 0
emp_title : 0

emp_length : 0
home_ownership : 0
annual_inc : 0
verification_status : 0

issue_d : 0
loan_status : 0
purpose : 0
dti : 0

earliest_cr_line : 0
open_acc : 0
pub_rec : 0
revol_bal : 0
revol_util : 0
total_acc : 0
initial_list_status : 0
application_type : 0
mort_acc : 0
pub_rec_bankruptcies : 0

city : 0

state : 0

zipcode : 0
issue_year : 0

earliest_cr_line_year : 0
issue_month : 0
earliest_cr_line_month : 0

```
In [ ]: # save to csv
df.write.csv("logistic_regression_cleaned.csv", header=True, mode='overwrite')
# save as parquet
df.write.parquet("logistic_regression_cleaned.parquet", mode='overwrite')
```

EDA

```
In [ ]: df = spark.read \
        .option("header", "true") \
        .option("inferSchema", "true") \
        .option("multiline", "true") \
        .option("escape", "\\\"") \
        .csv("./logistic_regression_cleaned.csv")
df.cache();
```

24/07/31 18:50:43 WARN CacheManager: Asked to cache already cached data.

```
In [ ]: df.createOrReplaceTempView("data")
pdf = df.toPandas();
```

```
In [ ]: pdf.head().T
```

	0	1	2	3	4
grade	B	B	B	A	C
loan_amnt	10000.0	8000.0	15600.0	7200.0	24375.0
term	36	36	36	36	60
int_rate	11.44	11.99	10.49	6.49	17.27
installment	329.48	265.68	506.97	220.65	609.33
sub_grade	B4	B5	B3	A2	C5
emp_title	marketing	credit analyst	statistician	client advocate	destiny management inc.
emp_length	10	4	1	6	9
home_ownership	RENT	MORTGAGE	RENT	RENT	MORTGAGE
annual_inc	117000	65000	43057	54000	55000
verification_status	Not Verified	Not Verified	Source Verified	Not Verified	Verified
issue_d	2015-01-01	2015-01-01	2015-01-01	2014-11-01	2013-04-01
loan_status	Fully Paid	Fully Paid	Fully Paid	Fully Paid	Charged Off
purpose	vacation	debt_consolidation	credit_card	credit_card	credit_card
dti	26.24	22.05	12.79	2.6	33.95
earliest_cr_line	2015-01-01	2015-01-01	2015-01-01	2014-11-01	2013-04-01
open_acc	16	17	13	6	13
pub_rec	0	0	0	0	0
revol_bal	36369	20131	11987	5472	24584
revol_util	41.8	53.3	92.2	21.5	69.8
total_acc	25	27	26	13	43
initial_list_status	w	f	f	f	f
application_type	INDIVIDUAL	INDIVIDUAL	INDIVIDUAL	INDIVIDUAL	INDIVIDUAL
mort_acc	0	3	0	0	1
pub_rec_bankruptcies	0	0	0	0	0
city	Mendozaberg	Loganmouth	New Sabrina	Delacruzside	Greggshire
state	OK	SD	WV	MA	VA
zipcode	22690	5113	5113	813	11650
issue_year	2015	2015	2015	2014	2013
earliest_cr_line_year	2015	2015	2015	2014	2013
issue_month	1	1	1	11	4
earliest_cr_line_month	1	1	1	11	4

```
In [ ]: pdf.describe().T
```

	count	mean	std	min	25%	50%	75%	max
loan_amnt	396030.0	14113.888089	8357.441341	500.00	8000.00	12000.00	20000.00	40000.00
term	396030.0	41.698053	10.212038	36.00	36.00	36.00	36.00	60.00
int_rate	396030.0	13.639400	4.472157	5.32	10.49	13.33	16.49	30.99
installment	396030.0	431.849698	250.727790	16.08	250.33	375.43	567.30	1533.81
emp_length	396030.0	6.006532	3.437114	1.00	3.00	6.00	10.00	10.00
annual_inc	396030.0	74203.170926	61637.622333	0.00	45000.00	64000.00	90000.00	8706582.00
dti	396030.0	17.379514	18.019092	0.00	11.28	16.91	22.98	9999.00
open_acc	396030.0	11.311153	5.137649	0.00	8.00	10.00	14.00	90.00
pub_rec	396030.0	0.178191	0.530671	0.00	0.00	0.00	0.00	86.00
revol_bal	396030.0	15844.539853	20591.836109	0.00	6025.00	11181.00	19620.00	1743266.00
revol_util	396030.0	53.754260	24.484857	0.00	35.80	54.80	72.90	892.30
total_acc	396030.0	25.414744	11.886991	2.00	17.00	24.00	32.00	151.00
mort_acc	396030.0	1.682895	2.087995	0.00	0.00	1.00	3.00	34.00
pub_rec_bankruptcies	396030.0	0.121483	0.355962	0.00	0.00	0.00	0.00	8.00
zipcode	396030.0	33998.447686	25605.865779	813.00	11650.00	29597.00	48052.00	93700.00
issue_year	396030.0	2013.629074	1.481725	2007.00	2013.00	2014.00	2015.00	2016.00
earliest_cr_line_year	396030.0	2013.629074	1.481725	2007.00	2013.00	2014.00	2015.00	2016.00
issue_month	396030.0	6.553188	3.426622	1.00	4.00	7.00	10.00	12.00
earliest_cr_line_month	396030.0	6.553188	3.426622	1.00	4.00	7.00	10.00	12.00


```
In [ ]: pdf.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 396030 entries, 0 to 396029
Data columns (total 32 columns):
#   Column                                Non-Null Count  Dtype
---  ---                                -
0   grade                                396030 non-null  object
1   loan_amnt                            396030 non-null  float64
2   term                                396030 non-null  int32
3   int_rate                             396030 non-null  float64
4   installment                          396030 non-null  float64
5   sub_grade                           396030 non-null  object
6   emp_title                           396030 non-null  object
7   emp_length                          396030 non-null  int32
8   home_ownership                      396030 non-null  object
9   annual_inc                          396030 non-null  int32
10  verification_status                 396030 non-null  object
11  issue_d                             396030 non-null  object
12  loan_status                         396030 non-null  object
13  purpose                             396030 non-null  object
14  dti                                 396030 non-null  float64
15  earliest_cr_line                    396030 non-null  object
16  open_acc                            396030 non-null  int32
17  pub_rec                             396030 non-null  int32
18  revol_bal                           396030 non-null  int32
19  revol_util                          396030 non-null  float64
20  total_acc                           396030 non-null  int32
21  initial_list_status                 396030 non-null  object
22  application_type                    396030 non-null  object
23  mort_acc                           396030 non-null  int32
24  pub_rec_bankruptcies                396030 non-null  int32
25  city                                396030 non-null  object
26  state                                396030 non-null  object
27  zipcode                             396030 non-null  int32
28  issue_year                          396030 non-null  int32
29  earliest_cr_line_year               396030 non-null  int32
30  issue_month                         396030 non-null  int32
31  earliest_cr_line_month              396030 non-null  int32
dtypes: float64(5), int32(14), object(13)
memory usage: 75.5+ MB
```

```
In [ ]: cat_cols = ['term', 'grade', 'emp_title', 'emp_length', 'home_ownership', 'verification_status', 'application_type', 'purpose', 'city', 'state', 'zipcode', 'issue_month', 'issue_year', 'initial_list_status', 'mort_acc', 'pub_rec_bankruptcies', 'city', 'state', 'zipcode', 'issue_year', 'issue_month', 'earliest_cr_line_year', 'earliest_cr_line_month']
int_columns = pdf.select_dtypes(include=['int64', 'int32', "float64"]).columns
```

Outlier Check

```
In [ ]: fig, ax = plt.subplots(1, 5, figsize=(20, 5))
sns.boxplot(y='annual_inc', data=pdf, ax=ax[0]);
sns.boxplot(y='int_rate', data=pdf, ax=ax[1]);
sns.boxplot(y='loan_amnt', data=pdf, ax=ax[2]);
sns.boxplot(y='installment', data=pdf, ax=ax[3]);
sns.boxplot(y='total_acc', data=pdf, ax=ax[4]);
```




Observations

- Annual Income has large number of outliers. This needs to be handled while preparing model

```
In [ ]: c=['pub_rec_bankruptcies', 'dti', 'open_acc', 'revol_bal', 'revol_util', 'pub_rec']
fig, ax = plt.subplots(1, 6, figsize=(20, 5));

for i in range(6):
    sns.boxplot(y=c[i], data=pdf, ax=ax[i]);
plt.show();
```




Observations

- Revolve balance have higher number of outliers

Normality Check


```
In [ ]: fig, ax = plt.subplots(1, 4, figsize=(20, 5))
sns.kdeplot(x='loan_amnt', data=pdf, ax=ax[0]);
sns.kdeplot(x='int_rate', data=pdf, ax=ax[1]);
sns.kdeplot(x='annual_inc', data=pdf, ax=ax[2]);
sns.kdeplot(x='installment', data=pdf, ax=ax[3]);
```



Observations

- Above plots show that most of features have right skewed distribution


```
In [ ]: fig, ax = plt.subplots(1, 4, figsize=(20, 5))
sns.kdeplot(x='dti', data=pdf, ax=ax[0]);
sns.kdeplot(x='open_acc', data=pdf, ax=ax[1]);
sns.kdeplot(x='revol_bal', data=pdf, ax=ax[2]);
sns.kdeplot(x='total_acc', data=pdf, ax=ax[3]);
```



Observations

We can see that annual income, revol_balance and dti is highly right skewed.

```
In [ ]: fig, ax = plt.subplots(1, 4, figsize=(20, 5))
sns.countplot(x='emp_length', data=pdf, ax=ax[0]);
sns.countplot(x='term', data=pdf, ax=ax[1]);
sns.countplot(x='grade', data=pdf, ax=ax[2]);
sns.countplot(x='loan_status', data=pdf, ax=ax[3]);
```




Observations

- From above plot we can see that most loans are taken by users who are younger and decreases as experience increases
- Most of the loans are taken for 36 months
- Most of the loans are taken by B grade users

Bivariate Analysis


```
In [ ]: fig, ax = plt.subplots(1, 5, figsize=(20, 5))
sns.boxplot(x='loan_status', y='loan_amnt', data=pdf, ax=ax[0]);
sns.boxplot(x='loan_status', y='annual_inc', data=pdf[pdf["annual_inc"]<100000], ax=ax[1]); #adjusting for outliers
sns.boxplot(x='loan_status', y='emp_length', data=pdf, ax=ax[2]);
sns.boxplot(x='loan_status', y='int_rate', data=pdf, ax=ax[3]);
sns.boxplot(x='loan_status', y='total_acc', data=pdf, ax=ax[4]);
```

Observations

- From above plot we can see that int_rate can be a useful feature when differentiating users based on loan status


```
In [ ]: fig, ax = plt.subplots(1, 4, figsize=(20, 5))
sns.boxplot(x='loan_status', y='installment', data=pdf, ax=ax[0]);
sns.boxplot(x='loan_status', y='mort_acc', data=pdf, ax=ax[1]);
sns.boxplot(x='loan_status', y='dti', data=pdf[pdf["dti"]<100], ax=ax[2]);
sns.boxplot(x='loan_status', y='open_acc', data=pdf, ax=ax[3]);
```



Observations

- From above plot we can see that there is not much relation between defaulters and Paid users with above features


```
In [ ]: fig, ax = plt.subplots(1, 4, figsize=(20, 5))
pd.crosstab(pdf['term'], pdf['loan_status'], normalize="index").plot(kind='bar', stacked=True, ax=ax[0]);
pd.crosstab(pdf['grade'], pdf['loan_status'], normalize="index").plot(kind='bar', stacked=True, ax=ax[1]);
pd.crosstab(pdf['home_ownership'], pdf['loan_status'], normalize="index").plot(kind='bar', stacked=True, ax=ax[2]);
pd.crosstab(pdf['verification_status'], pdf['loan_status'], normalize="index").plot(kind='bar', stacked=True, ax=ax[3]);
```



Observations

- We can see that more loans with 60 month repayment terms are "Charged off"
- From above plot we can see that Grade A users are those users who have most number of paid loans as compared to Grade G

```
In [ ]: fig, ax = plt.subplots(1, 4, figsize=(20, 5))
pd.crosstab(pdf['emp_length'], pdf['loan_status'], normalize="index").plot(kind='bar', stacked=True, ax=ax[0]);
pd.crosstab(pdf['purpose'], pdf['loan_status'], normalize="index").plot(kind='bar', stacked=True, ax=ax[1]);
pd.crosstab(pdf['application_type'], pdf['loan_status'], normalize="index").plot(kind='bar', stacked=True, ax=ax[2]);
pd.crosstab(pdf['zipcode'], pdf['loan_status'], normalize="index").plot(kind='bar', stacked=True, ax=ax[3]);
```




Observations

- We can see that there is equal proportion of loan paid off and charged off for employee experience
- We can see small business have highest percentage of "Charged Off" loans
- Above graph shows that zipcode can be an important feature in separating users on basis of loan status.
- Direct pay application type has higher number of defaulters as compared to joint application type

```
In [ ]: pd.crosstab(pdf['state'], pdf['loan_status'], normalize="index").plot(kind='bar', stacked=True, figsize=(20, 5));
```


<Figure size 2000x500 with 0 Axes>



Observations

- Above plots show that all states have equal proportion of paid and charged off loans

```
In [ ]: fig, ax = plt.subplots(1, 4, figsize=(20, 5))
sns.boxplot(x='loan_status', y='annual_inc', data=pdf[pdf["annual_inc"]<100000], ax=ax[0]); #adjusting for outliers
sns.boxplot(x='grade', y='int_rate', data=pdf, ax=ax[1], order=sorted(pdf['grade'].unique()));
sns.scatterplot(x='installment', y='loan_amnt', data=pdf, ax=ax[2]);
sns.histplot(x='loan_amnt', data=pdf, bins=50, hue='loan_status', ax=ax[3]);
```




Observations

- From above plot we can see that annual income is not important feature for loan status
- The loan amount is directly proportional to installment
- We can see that interest rate is directly proportional to grade

```
In [ ]: fig, ax = plt.subplots(1, 4, figsize=(20, 5))
spark.sql("""
select round(avg(loan_amnt),2) avg_loan, purpose from data group by purpose order by avg_loan desc
""").toPandas().plot(kind='bar', x='purpose', y='avg_loan', ax=ax[0]);
spark.sql("""
select round(avg(loan_amnt),2) avg_loan, grade from data group by grade order by avg_loan desc
""").toPandas().plot(kind='bar', x='grade', y='avg_loan', ax=ax[1]);


spark.sql("""
select round(avg(loan_amnt),2) avg_loan, application_type from data group by application_type order by avg_loan desc
""").toPandas().plot(kind='bar', x='application_type', y='avg_loan', ax=ax[2]);
spark.sql("""
select round(avg(loan_amnt),2) avg_loan, term from data group by term order by avg_loan desc
""").toPandas().plot(kind='bar', x='term', y='avg_loan', ax=ax[3]);
```



Observations

- Above plots show average loan amount with different features
- We can see that average loan amount for users is highest for houses and lowest for vacations
- The users with Joint application have higher loan amount
- G graded employees take the highest loan amount while B takes lowest


```
In [ ]: fig , ax = plt.subplots(1, 4, figsize=(20, 5))
spark.sql("""
select round(avg(loan_amnt),2) avg_loan, home_ownership  from data group by home_ownership order by avg_loan desc
""").toPandas().plot(kind='bar', x='home_ownership', y='avg_loan', ax=ax[0]);
spark.sql("""
select round(avg(loan_amnt),2) avg_loan, initial_list_status  from data group by initial_list_status order by avg_loan desc
""").toPandas().plot(kind='bar', x='initial_list_status', y='avg_loan', ax=ax[1]);
sns.scatterplot(x="pub_rec", y="loan_status", data=pdf, ax=ax[2]);
sns.scatterplot(x="pub_rec_bankruptcies", y="loan_status", data=pdf, ax=ax[3]);
```



Observations

- above plots shows the average loan amount for home ownership and initial list status

```
In [ ]: pdf["emp_title"].value_counts().head(51).iloc[1:].plot(kind='bar', figsize=(20, 5));
```




Observations

- Above plot shows top 50 employee titles with highest number of loans
- Above plot shows that employee with Manager and Teacher title, has the highest number of loans

```
In [ ]: spark.sql("""
select
  emp_title,
  -- sum(case when loan_status = 'Fully Paid' then 1 else 0 end)/count(*) as fully_paid,
  round(sum(case when loan_status = 'Charged Off' then 1 else 0 end)/count(*),2) as charged_off
from data
where emp_title != 'unknown' and loan_amnt > 10000
group by emp_title
having count(*) > 100
```



```
order by charged_off desc
""").toPandas().head(51).iloc[1:].plot(kind='bar', x='emp_title', y='charged_off', figsize=(20, 5));
```



Observations

Above is a list of employee titles who have higher ratio of charged off loans and have taken more than 100 loans above 10000.


```
In [ ]: pdf["loan_status"].value_counts(normalize=True)
```

```
Out [ ]: loan_status
Fully Paid    0.803871
Charged Off   0.196129
Name: proportion, dtype: float64
```

Observations

We can see that 19% users did not pay loans and 20% paid loans

```
In [ ]: cols=pdf.select_dtypes(include=['int64', 'int32', "float64"]).columns
corr = pdf[cols].corr()
plt.figure(figsize=(15, 13))
sns.heatmap(corr, annot=True, fmt=".2f", linewidths=0.5, cmap='viridis');
```



Observations

- Above plot shows correlation between different columns. We can see that loan_amnt and installment are highly correlated. Similarly, pub_rec and pub_rec_bankruptcies are highly correlated.

Model Building

```
In [ ]: df_model = spark.read.parquet("logistic_regression_cleaned.parquet")
df_model.cache();
# https://saturncloud.io/blog/feature-selection-in-pyspark-a-comprehensive-guide-for-data-scientists/
```

```
In [ ]: df_model.limit(5).toPandas()
```

	grade	loan_amnt	term	int_rate	installment	sub_grade	emp_title	emp_length	home_ownership	annual_inc	verification_status	issue_d	loan_status	purpose	dti	earliest_cr_line	open_acc
0	B	10000.0	36	11.44	329.48	B4	marketing	10	RENT	117000	Not Verified	2015-01-01	Fully Paid	vacation	26.24	2015-01-01	16
1	B	8000.0	36	11.99	265.68	B5	credit analyst	4	MORTGAGE	65000	Not Verified	2015-01-01	Fully Paid	debt_consolidation	22.05	2015-01-01	17
2	B	15600.0	36	10.49	506.97	B3	statistician	1	RENT	43057	Source Verified	2015-01-01	Fully Paid	credit_card	12.79	2015-01-01	13
3	A	7200.0	36	6.49	220.65	A2	client advocate	6	RENT	54000	Not Verified	2014-11-01	Fully Paid	credit_card	2.60	2014-11-01	6
4	C	24375.0	60	17.27	609.33	C5	destiny management inc.	9	MORTGAGE	55000	Verified	2013-04-01	Charged Off	credit_card	33.95	2013-04-01	13

Outlier Treatment

```
In [ ]: df_model = df_model.filter("dti < 50")
df_model = df_model.withColumn("revol_bal", sf.power(df_model["revol_bal"], 1/3))
df_model = df_model.withColumn("annual_inc", sf.power(df_model["annual_inc"], 1/3))
# df_model = df_model.withColumn("total_acc", sf.log(df_model["total_acc"]))
df_model = df_model.filter("revol_util < 200")
df_model = df_model.withColumn("loan_amnt", sf.log(df_model["loan_amnt"]))
# df_model = df_model.withColumn("installment", sf.log(df_model["installment"]))
# df_model = df_model.withColumn("open_acc", sf.log1p(df_model["open_acc"]))
```

Drop Columns

```
In [ ]: drop_cols=['issue_d', 'emp_title', 'earliest_cr_line', 'emp_title', 'grade','city' ]
df_model = df_model.drop(*drop_cols)
```

Encodings

```
In [ ]: sub_grade = df_model.select("sub_grade").distinct().sort("sub_grade")
indexer = StringIndexer(inputCol="sub_grade", outputCol="sub_grade_index")
indexer_model = indexer.fit(sub_grade)
df_encoded = indexer_model.transform(sub_grade)
df_model = indexer_model.transform(df_model)
df_model=df_model.drop("sub_grade")
df_model = df_model.withColumn("sub_grade", df_model["sub_grade_index"].cast(IntegerType()))
df_model = df_model.drop("sub_grade_index")
```

```
In [ ]: df_model = df_model.withColumn("loan_status", sf.when(df_model["loan_status"] == "Fully Paid", 0).otherwise(1).cast(IntegerType()))
df_model = df_model.withColumn("term", sf.when(df_model["term"] == 36, 0).otherwise(1).cast(IntegerType()))
# df_model = df_model.withColumn("application_type", sf.when(df_model["application_type"] == "INDIVIDUAL", 0).otherwise(1).cast(IntegerType()))
df_model = df_model.withColumn("initial_list_status", sf.when(df_model["initial_list_status"] == "w", 0).otherwise(1).cast(IntegerType()))
```

```
In [ ]: def target_mean_encoding(df, col, target):
    """
    :param df: pyspark.sql.dataframe
        dataframe to apply target mean encoding
    :param col: str list
        list of columns to apply target encoding
    :param target: str
        target column
    :return:
        dataframe with target encoded columns
    """
    target_encoded_columns_list = []
    for c in col:
        means = df.groupby(sf.col(c)).agg(sf.mean(target).alias(f"{c}_mean_encoding"))
        dict_ = means.toPandas().to_dict()
        target_encoded_columns = [sf.when(sf.col(c) == v, encoder)
                                   for v, encoder in zip(dict_[c].values(),dict_[f"{c}_mean_encoding"].values())]
        target_encoded_columns_list.append(sf.coalesce(*target_encoded_columns).alias(f"{c}_mean_encoding"))
    return df.select(*col, *target_encoded_columns_list).distinct()
```

```
In [ ]: df_target_encoded = target_mean_encoding(df_model, col=['state'], target='loan_status')
df_model = df_model.join(df_target_encoded, "state")
df_model = df_model.drop("state")
df_model = df_model.withColumnRenamed("state_mean_encoding", "state")
```

```
In [ ]: df_model = df_model.withColumn("issue_month_sin", sf.sin((df_model.issue_month-1)*(2.*np.pi/12)))
df_model = df_model.withColumn("issue_month_cos", sf.cos((df_model.issue_month-1)*(2.*np.pi/12)))
df_model = df_model.withColumn("earliest_cr_line_month_sin", sf.sin((df_model.earliest_cr_line_month-1)*(2.*np.pi/12)))
df_model = df_model.withColumn("earliest_cr_line_month_cos", sf.cos((df_model.earliest_cr_line_month-1)*(2.*np.pi/12)))
df_model = df_model.drop("issue_month", "earliest_cr_line_month")
```

```
In [ ]: df_model.limit(5).toPandas()
```

	loan_amnt	term	int_rate	installment	emp_length	home_ownership	annual_inc	verification_status	loan_status	purpose	dti	open_acc	pub_rec	revol_bal	revol_util	total_acc	initial_list_stat
0	9.210340	0	11.44	329.48	10	RENT	48.909732	Not Verified	0	vacation	26.24	16	0	33.131705	41.8	25	
1	8.987197	0	11.99	265.68	4	MORTGAGE	40.207258	Not Verified	0	debt_consolidation	22.05	17	0	27.203312	53.3	27	
2	9.655026	0	10.49	506.97	1	RENT	35.049454	Source Verified	0	credit_card	12.79	13	0	22.886014	92.2	26	
3	8.881836	0	6.49	220.65	6	RENT	37.797631	Not Verified	0	credit_card	2.60	6	0	17.621736	21.5	13	
4	10.101313	1	17.27	609.33	9	MORTGAGE	38.029525	Verified	1	credit_card	33.95	13	0	29.077084	69.8	43	

```
In [ ]: df_model.write.parquet("df_model.parquet", mode='overwrite')
```

```
In [ ]: df_model = pd.read_parquet("df_model.parquet")
df_model.head()
```

	loan_amnt	term	int_rate	installment	emp_length	home_ownership	annual_inc	verification_status	loan_status	purpose	dti	open_acc	pub_rec	revol_bal	revol_util	total_acc	initial_list_stat
0	9.210340	0	11.44	329.48	10	RENT	48.909732	Not Verified	0	vacation	26.24	16	0	33.131705	41.8	25	
1	8.987197	0	11.99	265.68	4	MORTGAGE	40.207258	Not Verified	0	debt_consolidation	22.05	17	0	27.203312	53.3	27	
2	9.655026	0	10.49	506.97	1	RENT	35.049454	Source Verified	0	credit_card	12.79	13	0	22.886014	92.2	26	
3	8.881836	0	6.49	220.65	6	RENT	37.797631	Not Verified	0	credit_card	2.60	6	0	17.621736	21.5	13	
4	10.101313	1	17.27	609.33	9	MORTGAGE	38.029525	Verified	1	credit_card	33.95	13	0	29.077084	69.8	43	

```
In [ ]: df_model["was_bankrupt"] = df_model["pub_rec_bankruptcies"].apply(lambda x: 1 if x > 0 else 0)
df_model["has_pub_rec"] = df_model["pub_rec"].apply(lambda x: 1 if x > 0 else 0)
df_model["has_mort_acc"] = df_model["mort_acc"].apply(lambda x: 1 if x > 0 else 0)
df_model = df_model.drop(["pub_rec_bankruptcies", "pub_rec", "mort_acc"], axis=1)
```

```
In [ ]: df_model.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 395994 entries, 0 to 395993
Data columns (total 29 columns):
#   Column                Non-Null Count  Dtype
---  -
0   loan_amnt              395994 non-null float64
1   term                  395994 non-null int32
2   int_rate              395994 non-null float64
3   installment           395994 non-null float64
4   emp_length            395994 non-null int32
5   home_ownership         395994 non-null object
6   annual_inc            395994 non-null float64
7   verification_status   395994 non-null object
8   loan_status           395994 non-null int32
9   purpose               395994 non-null object
10  dti                   395994 non-null float64
11  open_acc              395994 non-null int32
12  revol_bal             395994 non-null float64
13  revol_util            395994 non-null float64
14  total_acc             395994 non-null int32
15  initial_list_status   395994 non-null int32
16  application_type      395994 non-null object
17  zipcode               395994 non-null object
18  issue_year            395994 non-null int32
19  earliest_cr_line_year 395994 non-null int32
20  sub_grade             395994 non-null int32
21  state                 395994 non-null float64
22  issue_month_sin       395994 non-null float64
23  issue_month_cos       395994 non-null float64
24  earliest_cr_line_month_sin 395994 non-null float64
25  earliest_cr_line_month_cos 395994 non-null float64
26  was_bankrupt          395994 non-null int64
27  has_pub_rec           395994 non-null int64
28  has_mort_acc          395994 non-null int64
dtypes: float64(12), int32(9), int64(3), object(5)
memory usage: 74.0+ MB
```

```
In [ ]: df_model["loan_status"].value_counts(normalize=True)
```

```
Out[ ]: loan_status
0      0.803876
1      0.196124
Name: proportion, dtype: float64
```

```
In [ ]: oneHotEncoder = OneHotEncoder()
oneHotEncoder.fit(df_model[["home_ownership", "verification_status", "purpose", "zipcode", "application_type"]])
df_encoded = oneHotEncoder.transform(df_model[["home_ownership", "verification_status", "purpose", "zipcode", "application_type"]])
df_encoded_dataframe = pd.DataFrame(df_encoded.toarray(), columns=oneHotEncoder.get_feature_names_out(["home_ownership", "verification_status", "purpose", "zipcode", "application_type"]))
```

Out[]:


▼ OneHotEncoder ⓘ ⓘ

OneHotEncoder()

Finding Correlation

```
In [ ]: cols=df_model.select_dtypes(include=['int64', 'int32', "float64"]).columns
corr = df_model[cols].corr()
```

```
In [ ]: plt.figure(figsize=(22, 18))
sns.heatmap(corr, annot=True, fmt=".2f", cmap='viridis');
```



Observations

From above plot we can see that following features are correlated

- Interest Rate and SubGrade
- loan amount and Installement amount
- Earliest credit month sin and issue month sin are correlated

Dropping Correlated Features

```
In [ ]: to_drop = ["loan_amnt", "int_rate", "issue_month_cos", "issue_month_sin", "issue_year"]
df_model = df_model.drop(to_drop, axis=1)

In [ ]: df_model = pd.concat([df_model, df_encoded_dataframe], axis=1)
df_model = df_model.drop(["home_ownership", "verification_status", "purpose", "zipcode", "application_type"], axis=1)

In [ ]: df_model

Out [ ]:
      term  installment  emp_length  annual_inc  loan_status  dti  open_acc  revol_bal  revol_util  total_acc  initial_list_status  earliest_cr_line_year  sub_grade  state  earliest_cr_line_month_sin  earlie
0      0      329.48      10  48.909732      0  26.24      16  33.131705      41.8      25      0      2015      8  0.192013      0.000000
1      0      265.68      4  40.207258      0  22.05      17  27.203312      53.3      27      1      2015      9  0.197067      0.000000
2      0      506.97      1  35.049454      0  12.79      13  22.886014      92.2      26      1      2015      7  0.203976      0.000000
3      0      220.65      6  37.797631      0   2.60      6  17.621736      21.5      13      1      2014      1  0.195983     -0.866025
4      1      609.33      9  38.029525      1  33.95      13  29.077084      69.8      43      1      2013     14  0.195129      1.000000
...    ...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...
395989  1      217.38      2  34.199519      0  15.63      6  12.578177      34.3      23      0      2015      8  0.195557     -1.000000
395990  0      700.42      5  47.914199      0  21.45      6  35.105261      95.7      8      1      2015     10  0.193690      0.500000
395991  0      161.32     10  38.372151      0  17.56     15  31.979153      66.9     23      1      2013      5  0.186492     -1.000000
395992  1      503.02     10  40.000000      0  15.88      9  25.042063      53.8     20      1      2012     11  0.193469     -0.500000
395993  0       67.98     10  35.032894      0   8.32      3  16.251243      91.3     19      1      2010     11  0.198048      0.500000
```

395994 rows x 55 columns

```
In [ ]: X_train, X_test, y_train, y_test = train_test_split(df_model.drop("loan_status", axis=1), df_model["loan_status"], test_size=0.2, random_state=25)

In [ ]: min_max_scale_cols = ['earliest_cr_line_year', 'emp_length', 'sub_grade']
standard_scale_cols = ['annual_inc', 'dti', 'installment', 'open_acc', 'revol_bal', 'revol_util', 'total_acc',]
```



```
In [ ]: def perform_scaling(scaler, X_train, X_test, cols):
        X_train_scaled = X_train.copy()
        X_test_scaled = X_test.copy()
        X_train_scaled[cols] = scaler.fit_transform(X_train[cols])
        X_test_scaled[cols] = scaler.transform(X_test[cols])
        return X_train_scaled, X_test_scaled
```

```
In [ ]: X_train, X_test = perform_scaling(MinMaxScaler(), X_train, X_test, min_max_scale_cols)
        X_train, X_test = perform_scaling(StandardScaler(), X_train, X_test, standard_scale_cols)
```

```
In [ ]: X_train.head()
```

Out []:

	term	installment	emp_length	annual_inc	dti	open_acc	revol_bal	revol_util	total_acc	initial_list_status	earliest_cr_line_year	sub_grade	state	earliest_cr_line_month_sin	earliest_cr_li
	164395	1	-0.435095	1.000000	-0.292487	1.438927	-0.838918	1.122573	1.324278	-1.464612	1	0.666667	0.441176	0.198023	-0.500000
	241040	0	-0.917292	0.000000	-0.940833	-0.861258	-1.227969	-0.836158	1.557287	-0.959992	0	0.777778	0.323529	0.194740	-0.866025
	365230	0	-0.430867	0.777778	-0.202413	-0.110529	-0.644393	-0.611847	-0.654258	-1.380508	1	0.555556	0.058824	0.195210	-0.866025
	26021	0	0.298293	0.000000	-0.109585	-0.542506	-0.255341	-0.937393	-1.161155	0.974386	1	0.555556	0.264706	0.195557	0.866025
	305341	0	-0.925031	0.555556	-0.202413	-1.520915	-0.644393	-0.553120	0.429031	-1.296405	1	0.666667	0.470588	0.196346	0.866025

VIF

```
In [ ]: vif_cols = ['term', 'installment', 'emp_length', 'annual_inc', 'dti', 'open_acc',
                  'revol_bal', 'revol_util', 'total_acc', 'initial_list_status',
                  'earliest_cr_line_year', 'sub_grade',
                  'earliest_cr_line_month_sin', 'earliest_cr_line_month_cos', 'state',
                  'was_bankrupt', 'has_pub_rec', 'has_mort_acc']
```

```
In [ ]: X_vif = pd.DataFrame(X_train[vif_cols], columns=vif_cols)
        vif = pd.DataFrame()

        vif['Features'] = X_vif.columns
        vif['VIF'] = [variance_inflation_factor(X_vif.values, i) for i in range(X_vif.shape[1])]
        vif['VIF'] = round(vif['VIF'], 2)
        vif = vif.sort_values(by = "VIF", ascending = False)
        vif
```

Out []:

	Features	VIF
14	state	38.71
10	earliest_cr_line_year	28.37
11	sub_grade	6.26
16	has_pub_rec	4.86
15	was_bankrupt	4.64
2	emp_length	3.32
9	initial_list_status	3.14
17	has_mort_acc	3.11
5	open_acc	2.22
8	total_acc	2.16
6	revol_bal	2.06
0	term	1.84
3	annual_inc	1.80
7	revol_util	1.62
1	installment	1.52
4	dti	1.46
12	earliest_cr_line_month_sin	1.03
13	earliest_cr_line_month_cos	1.01

```
In [ ]: vif_cols =list(set(vif_cols) - set(["state", "earliest_cr_line_year"]))

        X_vif = pd.DataFrame(X_train[vif_cols], columns=vif_cols)
        vif = pd.DataFrame()

        vif['Features'] = X_vif.columns
        vif['VIF'] = [variance_inflation_factor(X_vif.values, i) for i in range(X_vif.shape[1])]
        vif['VIF'] = round(vif['VIF'], 2)
        vif = vif.sort_values(by = "VIF", ascending = False)
        vif
```

Out []:

	Features	VIF
5	has_pub_rec	4.83
13	was_bankrupt	4.64
9	sub_grade	4.18
8	emp_length	2.78
3	has_mort_acc	2.60
12	open_acc	2.21
4	initial_list_status	2.18
14	total_acc	2.14
7	revol_bal	2.05
0	term	1.81
15	annual_inc	1.79
11	revol_util	1.55
1	installment	1.49
6	dti	1.42
2	earliest_cr_line_month_cos	1.00
10	earliest_cr_line_month_sin	1.00

```
In [ ]: vif_cols
```



```
Out [ ]: ['term',
          'installment',
          'earliest_cr_line_month_cos',
          'has_mort_acc',
          'initial_list_status',
          'has_pub_rec',
          'dti',
          'revol_bal',
          'emp_length',
          'sub_grade',
          'earliest_cr_line_month_sin',
          'revol_util',
          'open_acc',
          'was_bankrupt',
          'total_acc',
          'annual_inc']
```

```
In [ ]: X_train = X_train.drop(["state", "earliest_cr_line_year"], axis=1)
X_test = X_test.drop(["state", "earliest_cr_line_year"], axis=1)
```

```
In [ ]: from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import GridSearchCV
from sklearn.feature_selection import RFE
from sklearn.metrics import roc_auc_score, roc_curve, confusion_matrix, classification_report
from sklearn.metrics import f1_score, precision_score, recall_score
from sklearn.metrics import precision_recall_curve, auc, ConfusionMatrixDisplay
import pickle
```

Hyperparameter Tuning

```
In [ ]: penalty = ['l1', 'l2']
lambdas = [0.01, 0.1, 1, 10, 100, 1000, 10000]
C_values = [1/lambda_val for lambda_val in lambdas]
param_grid = {'C': C_values, 'penalty': penalty}
model = LogisticRegression(max_iter=1000 , penalty=penalty, class_weight='balanced', solver='liblinear', n_jobs=-1)
```

```
In [ ]: grid_search = GridSearchCV(model, param_grid, scoring='f1', n_jobs=-1, cv=2, verbose=3)
grid_search.fit(X_train, y_train)
```

Fitting 2 folds for each of 14 candidates, totalling 28 fits

[CV 1/2] ENDC=10.0, penalty=l1;	score=0.625	total time=	12.8s
[CV 1/2] ENDC=100.0, penalty=l1;	score=0.624	total time=	13.0s
[CV 2/2] ENDC=100.0, penalty=l1;	score=0.622	total time=	13.0s
[CV 2/2] ENDC=10.0, penalty=l1;	score=0.622	total time=	13.1s
[CV 2/2] ENDC=0.1, penalty=l2;	score=0.622	total time=	18.9s
[CV 1/2] ENDC=0.1, penalty=l2;	score=0.625	total time=	19.9s
[CV 2/2] ENDC=1.0, penalty=l2;	score=0.622	total time=	23.5s
[CV 1/2] ENDC=1.0, penalty=l2;	score=0.625	total time=	24.3s
[CV 1/2] ENDC=0.01, penalty=l2;	score=0.625	total time=	12.7s
[CV 2/2] ENDC=0.01, penalty=l2;	score=0.622	total time=	13.0s
[CV 1/2] ENDC=0.0001, penalty=l1;	score=0.557	total time=	3.0s
[CV 2/2] ENDC=0.0001, penalty=l1;	score=0.559	total time=	3.2s
[CV 1/2] ENDC=0.0001, penalty=l1;	score=0.622	total time=	11.1s
[CV 1/2] ENDC=0.0001, penalty=l2;	score=0.622	total time=	6.9s
[CV 2/2] ENDC=10.0, penalty=l2;	score=0.622	total time=	31.4s
[CV 2/2] ENDC=100.0, penalty=l2;	score=0.622	total time=	31.5s
[CV 2/2] ENDC=0.0001, penalty=l2;	score=0.619	total time=	7.0s
[CV 2/2] ENDC=0.0001, penalty=l1;	score=0.619	total time=	11.5s
[CV 1/2] ENDC=0.0001, penalty=l2;	score=0.604	total time=	3.6s
[CV 2/2] ENDC=0.0001, penalty=l2;	score=0.600	total time=	3.5s
[CV 1/2] ENDC=10.0, penalty=l2;	score=0.625	total time=	33.9s
[CV 1/2] ENDC=100.0, penalty=l2;	score=0.625	total time=	34.9s
[CV 1/2] ENDC=0.01, penalty=l1;	score=0.624	total time=	25.7s
[CV 2/2] ENDC=0.01, penalty=l1;	score=0.622	total time=	26.1s
[CV 2/2] ENDC=1.0, penalty=l1;	score=0.622	total time=	51.2s
[CV 1/2] ENDC=0.1, penalty=l1;	score=0.625	total time=	1.1min
[CV 1/2] ENDC=1.0, penalty=l1;	score=0.625	total time=	1.1min
[CV 2/2] ENDC=0.1, penalty=l1;	score=0.622	total time=	1.2min

Out []:

GridSearchCV ⓘ ⓘ

▶ estimator: LogisticRegression

▶ LogisticRegression

```
In [ ]: grid_search.best_params_
```

```
Out [ ]: {'C': 0.1, 'penalty': 'l1'}
```

Training Model with Updated values

```
In [ ]: model = LogisticRegression(max_iter=1000 , C=0.1, penalty="l1", class_weight='balanced', solver='liblinear', n_jobs=-1)
model.fit(X_train, y_train)
```

Out []:

LogisticRegression ⓘ

LogisticRegression(C=0.1, class_weight='balanced', max_iter=1000, n_jobs=1, penalty='l1', solver='liblinear')

```
In [ ]: # with open("logistic_regression_model.pkl", "wb") as f:
#     pickle.dump(model, f)

with open("logistic_regression_model.pkl", "rb") as f:
    model = pickle.load(f)
```

```
In [ ]: y_test_pred = model.predict(X_test)
y_train_pred = model.predict(X_train)
y_test_prob = model.predict_proba(X_test)
y_train_prob = model.predict_proba(X_train)
```

Calculate best threshold based on f1 score

```
In [ ]: thresholds = np.arange(0.0, 1.0, 0.01)
```

```
precisions = []
recalls = []
f1s = []

for threshold in thresholds:
    _y_test_pred = (y_test_prob[:, 1] >= threshold).astype(int)
    precisions.append(precision_score(y_test, _y_test_pred))
    recalls.append(recall_score(y_test, _y_test_pred))
    f1s.append(f1_score(y_test, _y_test_pred))


best_threshold = thresholds[np.argmax(f1s)]

print(f'Best threshold: {best_threshold}')
print(f'Precision at best threshold: {precisions[np.argmax(f1s)]}')
print(f'Recall at best threshold: {recalls[np.argmax(f1s)]}')
print(f'F1 score at best threshold: {f1s[np.argmax(f1s)]}')
```


Best threshold: 0.67
Precision at best threshold: 0.7145569139646138
Recall at best threshold: 0.6050279688806018
F1 score at best threshold: 0.6552468491052155

```
In [ ]: # best_threshold = 0.66
y_test_pred = (y_test_prob[:, 1] >= best_threshold).astype(int)
y_train_pred = (y_train_prob[:, 1] >= best_threshold).astype(int)
```

```
In [ ]: coefficients = model.coef_[0]
feature_importance = pd.DataFrame({'Feature': X_train.columns, 'Importance': coefficients})
feature_importance = feature_importance.sort_values('Importance', ascending=True)
feature_importance.plot(x='Feature', y='Importance', kind='barh', figsize=(10, 15));
```



```
In [ ]: coefficients = model.coef_[0]
feature_importance = pd.DataFrame({'Feature': X_train.columns, 'Importance': np.abs(coefficients)})
feature_importance = feature_importance.sort_values('Importance', ascending=False)
feature_importance.plot(x='Feature', y='Importance', kind='barh', figsize=(10, 15));
```



Results Evaluation

ROC Curve


```
In [ ]: fpr, tpr, threshold = roc_curve(y_test, y_test_prob[:,1], pos_label=1)

random_probs = [0 for i in range(len(y_test))]
p_fpr, p_tpr, _ = roc_curve(y_test, random_probs, pos_label=1)

auc_score = roc_auc_score(y_test, y_test_prob[:,1])
print("AUC Score: ",auc_score)
```

AUC Score: 0.9055573414285613

```
In [ ]: plt.plot(fpr, tpr, linestyle='--',color='orange', label='Logistic Regression')
plt.plot(p_fpr, p_tpr, linestyle='--', color='blue')
plt.title('ROC curve')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive rate')
plt.show();
```




Observations

- From above plot we can see that there is good balance of TPR and FPR

Precision recall Curve

```
In [ ]: precision, recall, thr = precision_recall_curve(y_test, y_test_prob[:,1])
plt.plot(recall, precision)
plt.xlabel('Recall')
```

```
plt.ylabel('Precision')
plt.title('PR curve')
plt.show();
```



Observations

Above plot shows value of precision against recall

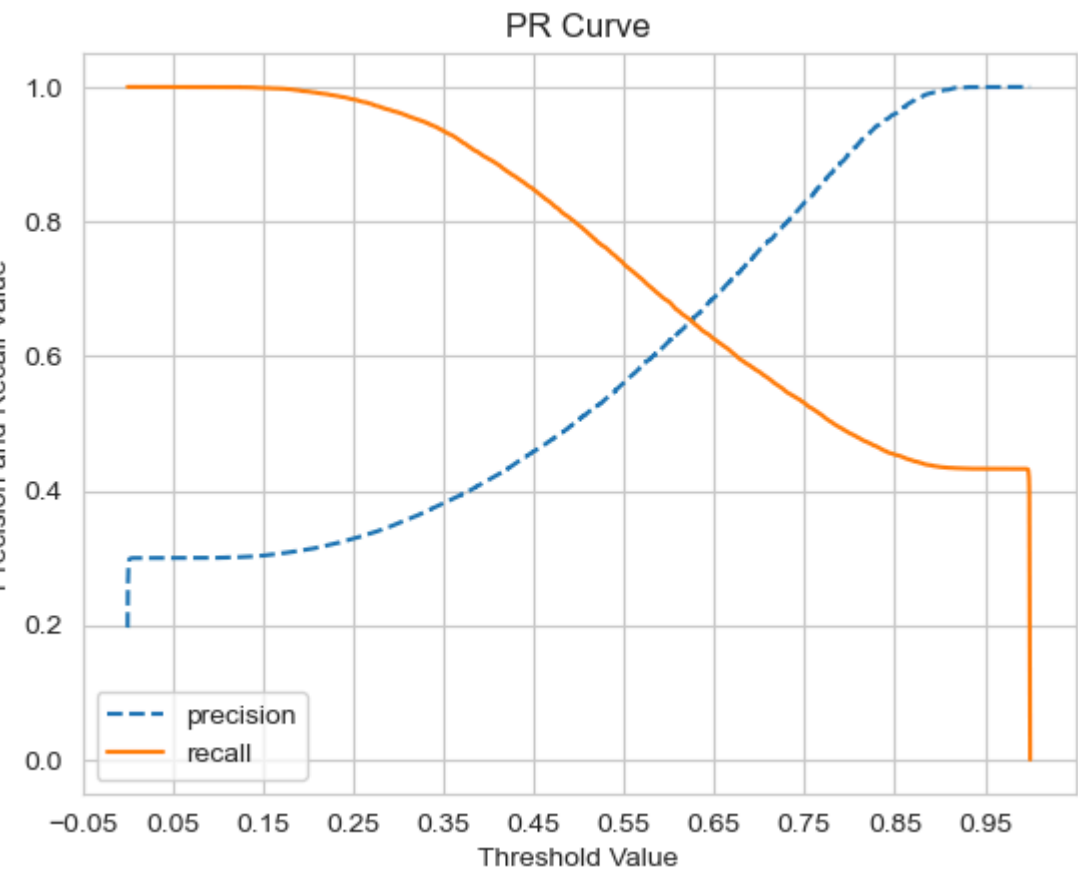
```
In [ ]: precision, recall, thr = precision_recall_curve(y_test, y_test_prob[:,1])

threshold_boundary = thr.shape[0]

plt.plot(thr, precision[0:threshold_boundary], linestyle='--', label='precision')
plt.plot(thr, recall[0:threshold_boundary], label='recall')

start, end = plt.xlim()
plt.xticks(np.round(np.arange(start, end, 0.1), 2))

plt.xlabel('Threshold Value')
plt.ylabel('Precision and Recall Value')
plt.legend()
plt.title("PR Curve")
plt.show();
```




Observations

- From above plot we can see that as threshold value increases, precesion value increases.
- At the same time the value of recall decreases

Confusion Matrix

```
In [ ]: cm = confusion_matrix(y_test, y_test_pred, normalize='true', labels=[1, 0])
sns.heatmap(cm, annot=True, cmap='viridis');
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual');
```



Observations

- Above is the confusion matrix for the model
- 94% of paid loans were correctly predicted
- 61% of the unpaid loans were correctly predicted
- 39% of unpaid were predicted as paid
- 5% of the paid loans were predicted as not paid

Test Classification Report

```
In [ ]: print(classification_report(y_test, y_test_pred))
```

	precision	recall	f1-score	support
0	0.91	0.94	0.92	63646
1	0.71	0.61	0.66	15553
accuracy			0.87	79199
macro avg	0.81	0.77	0.79	79199
weighted avg	0.87	0.87	0.87	79199

Train Classification Report

```
In [ ]: print(classification_report(y_train, y_train_pred))
```

	precision	recall	f1-score	support
0	0.91	0.94	0.92	254684
1	0.71	0.60	0.66	62111
accuracy			0.88	316795
macro avg	0.81	0.77	0.79	316795
weighted avg	0.87	0.88	0.87	316795

Observations

- Class 0 (Defaulted Loans)
 - Precision (0.91):** Out of all the loans predicted to be defaulted, 91% were actually defaulted. This means that the model is quite accurate in predicting defaults and has a low false positive rate for this class.
 - Recall (0.94):** Out of all the actual defaulted loans, 94% were correctly identified by the model. This indicates that the model is very effective at capturing most of the defaulted loans, with a low false negative rate.
 - F1-Score (0.92):** It shows that the model has a good balance between precision and recall for defaulted loans.
- Class 1 (Paid Off Loans)
 - Precision (0.71):** Out of all the loans predicted to be paid off, 71% were actually paid off. This means that the model has a higher false positive rate for this class compared to class 0.
 - Recall (0.60):** Out of all the actual paid off loans, 60% were correctly identified by the model. This indicates that the model misses a significant number of loans that were actually paid off (higher false negative rate).
 - F1-Score (0.66):** The F1-score for paid off loans is 0.66, indicating a moderate balance between precision and recall for this class, but not as strong as for defaulted loans.

```
In [ ]: scores = {
    "f1":{
        "test": f1_score(y_test, y_test_pred),
        "train": f1_score(y_train, y_train_pred)
    },
    "precision":{
        "test": precision_score(y_test, y_test_pred),
        "train": precision_score(y_train, y_train_pred)
    },
    "recall":{
        "test": recall_score(y_test, y_test_pred),
        "train": recall_score(y_train, y_train_pred)
    },
    "auc":{
        "test": roc_auc_score(y_test, y_test_prob[:,1]),
        "train": roc_auc_score(y_train, y_train_prob[:,1])
    }
}
pd.DataFrame(scores)
```

	f1	precision	recall	auc
test	0.655247	0.714557	0.605028	0.905557
train	0.655008	0.714704	0.604514	0.906478

Conclusion

Precision-Recall Tradeoff

Increasing Precision

Precision is the ratio of true positives (correctly predicted positive instances) to the total predicted positives (true positives + false positives). To increase precision, you need to decrease the number of false positives. In the context of loan predictions: A false positive occurs when the model predicts that a loan will be paid off (positive prediction) but it is not paid off Reducing these false positives (incorrectly predicted paid-off loans) will increase precision.

Increasing Recall

Recall is the ratio of true positives to the total actual positives (true positives + false negatives). To increase recall, you need to decrease the number of false negatives. In the context of loan predictions: A false negative occurs when the model predicts that a loan will not be paid off (negative prediction) but it is actually paid off (positive actual). Reducing these false negatives (missed paid-off loans) will increase recall.

To Summarize Precision: Focus on minimizing false positives (loans predicted to be paid off but actually defaulted). Recall: Focus on minimizing false negatives (loans predicted to default but actually paid off). Balancing precision and recall often involves trade-offs, and the choice of threshold can help you achieve the desired balance depending on your specific goals.

Main Metric

According to me the bank should to focus on f1 score as the main metric because the value of f1 score is derived from precision and recall. This will keep both loss due to bad loans and profit due to good loan high.

Predictive Modelling

Model Performance

- The above model has high precision when predicting loan defaulters, but has mediocre precision for predicting whether a loan will be paid or not.
- This model is better for minimizing risk, but will take a hit on maximizing profitability.
- This can be adjusted by changing the threshold value, but trying to increase of decrease precision might affect recall value.

Gap between Precision and Recall

- If the gap between precision and recall increases, this might lead to incorrect prediction. This might lead to losses in business.

Insights and Answers to Questionnaire

- Geographical features had big influence on the outcome
- Zipcode feature has the strongest influence on the model.
- Subgrade, application type and term is next 3 important feature.
- Around 80% of the users have paid their loans
- Loan and loan installment have 95% correlation
- People with Grade A have high chances of paying off their loans
- Income of the users have very high right skew. This shows how diverse is the user base
- Managers and Teachers have highest numbers of loans
- We can see that average loan amount for users is highest for houses and lowest for vacations.
- The users with Joint application have higher loan amount.
- G graded employees take the highest loan amount while B takes lowest.

- House and small business purpose have highest number of loans
- Most loans are taken by users who are younger and decreases as experience increases
- Most of the loans are taken for 36 months
- Most of the loans are taken by B grade users
- Majority of the user have mortgage type home ownership
- House and small business have higher average loan amount

Recommendation

- A more powerful model can be used inplace of Logistic regression
- Users with zipcodes 11650, 86630, 93700 must not be given loans
- Users with zipcodes 00813, 05113, 29597 can be given loans
- Users with higher grades have higher default ratio. Other than higher interest rate, additional checking can be used to give loans
- Provide more options for the term period so customers can make payments according to their incomes.
- Model should be continously trained and threshold should be adjusted according to banks profit and losses
- Bank should deeply check Small businesses loans as these have highest charged off rate wrt to other categories
- Bank should encourage more joint pay loans as compared to direct pay and individual type as joint application types have highher repayment rate